

T5: Linear methods and Perceptron

Fundamentos del Aprendizaje Automático

Curso 2025/2026

Structure

① Linear models

Binary

Multiclass

Parameter estimation

② Perceptron

Introduction

Training

Limitations

Multiclass Perceptron

③ Multi-layer Perceptron

Introduction

Structure

Training

Training dynamics and regularization

Problem statement

- So far:

Problem statement

- So far:

T2: Parametric models for modeling *likelihood*: $P(\mathbf{x}|\omega) \approx \hat{P}(\mathbf{x}|\omega, \theta)$

Problem statement

- So far:

T2: Parametric models for modeling *likelihood*: $P(\mathbf{x}|\omega) \approx \hat{P}(\mathbf{x}|\omega, \theta)$

T4: Nonparametric density estimator for approximating *likelihood*

Problem statement

- So far:

T2: Parametric models for modeling *likelihood*: $P(\mathbf{x}|\omega) \approx \hat{P}(\mathbf{x}|\omega, \theta)$

T4: Nonparametric density estimator for approximating *likelihood*

T4: Distance-based classification for *posterior* $\hat{P}(\omega|\mathbf{x})$

Problem statement

- So far:

T2: Parametric models for modeling *likelihood*: $P(\mathbf{x}|\omega) \approx \hat{P}(\mathbf{x}|\omega, \theta)$

T4: Nonparametric density estimator for approximating *likelihood*

T4: Distance-based classification for *posterior* $\hat{P}(\omega|\mathbf{x})$

- **T5:** Parametric model for *discriminant functions* $\Rightarrow g(\mathbf{x}; \theta)$

Problem statement

- So far:

T2: Parametric models for modeling *likelihood*: $P(\mathbf{x}|\omega) \approx \hat{P}(\mathbf{x}|\omega, \boldsymbol{\theta})$
T4: Nonparametric density estimator for approximating *likelihood*
T4: Distance-based classification for *posterior* $\hat{P}(\omega|\mathbf{x})$

- **T5:** Parametric model for discriminant functions $\Rightarrow g(\mathbf{x}; \boldsymbol{\theta})$
→ $\boldsymbol{\theta}$ ⇒ weights of the model

Problem statement

- So far:

T2: Parametric models for modeling *likelihood*: $P(\mathbf{x}|\omega) \approx \hat{P}(\mathbf{x}|\omega, \boldsymbol{\theta})$
T4: Nonparametric density estimator for approximating *likelihood*
T4: Distance-based classification for *posterior* $\hat{P}(\omega|\mathbf{x})$

- **T5:** Parametric model for *discriminant functions* $\Rightarrow g(\mathbf{x}; \boldsymbol{\theta})$
 - $\boldsymbol{\theta}$ ⇒ *weights* of the model
 - $J(\boldsymbol{\theta})$ ⇒ *criterion function* to estimate $\boldsymbol{\theta}$

Outline

1 Linear models

Binary

Multiclass

Parameter estimation

2 Perceptron

Introduction

Training

Limitations

Multiclass Perceptron

3 Multi-layer Perceptron

Introduction

Structure

Training

Training dynamics and regularization

Linear discriminant function

Linear discriminant function

- Simplest type of ML model

Linear discriminant function


- Simplest type of ML model
- **Assumption:** output is a linear combination of the input features

Linear discriminant function

- Simplest type of ML model
- Assumption: output is a linear combination of the input features
- Two-category case ($\mathcal{W} = \{\omega_1, \omega_2\}$):

$$g(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x} + \theta_0$$

$\mathbf{x} \in \mathbb{R}^d$ feature vector
 $\boldsymbol{\theta} \in \mathbb{R}^d$ weight vector
 $\theta_0 \in \mathbb{R}$ weight threshold



Decision surface


- Decision surface (\mathcal{H}): Given by equation $g(\mathbf{x}) = 0$

Decision surface

- Decision surface (\mathcal{H}): Given by equation $g(\mathbf{x}) = 0$
 - If $g(\mathbf{x}) > 0$: Region $\mathcal{R}_1 \Rightarrow \hat{\omega} = \omega_1$
 - If $g(\mathbf{x}) < 0$: Region $\mathcal{R}_2 \Rightarrow \hat{\omega} = \omega_2$

Decision surface

- Decision surface (\mathcal{H}): Given by equation $g(\mathbf{x}) = 0$
 - If $g(\mathbf{x}) > 0$: Region $\mathcal{R}_1 \Rightarrow \hat{\omega} = \omega_1$
 - If $g(\mathbf{x}) < 0$: Region $\mathcal{R}_2 \Rightarrow \hat{\omega} = \omega_2$
- Geometric interpretation:
 - Vector θ : normal to surface \mathcal{H} and defines its *orientation*
 - Bias θ_0 : Shifts surface \mathcal{H} along θ



Decision surface

Formalization

- Multiclass: Generalizes to **one** linear function **per label**:

$$g_k(\mathbf{x}) = \boldsymbol{\theta}_k^T \mathbf{x} + \theta_{0k}, \quad k = 1, 2, \dots, |\mathcal{W}|$$

Formalization

- Multiclass: Generalizes to one linear function per label:

$$g_k(\mathbf{x}) = \boldsymbol{\theta}_k^T \mathbf{x} + \theta_{0k}, \quad k = 1, 2, \dots, |\mathcal{W}|$$

- Decision rule:

$$\hat{\omega} = \arg \max_{k=1, \dots, |\mathcal{W}|} g_k(\mathbf{x}) = \arg \max_{k=1, \dots, |\mathcal{W}|} \left(\boldsymbol{\theta}_k^T \mathbf{x} + \theta_{0k} \right)$$

Formalization

- Multiclass: Generalizes to one linear function per label:

$$g_k(\mathbf{x}) = \boldsymbol{\theta}_k^T \mathbf{x} + \theta_{0k}, \quad k = 1, 2, \dots, |\mathcal{W}|$$

- Decision rule:

$$\hat{\omega} = \arg \max_{k=1, \dots, |\mathcal{W}|} g_k(\mathbf{x}) = \arg \max_{k=1, \dots, |\mathcal{W}|} \left(\boldsymbol{\theta}_k^T \mathbf{x} + \theta_{0k} \right)$$

- Decision surface(s):

→ Each tuple $(\boldsymbol{\theta}_k^T, \theta_{0k})$ defines a hyperplane

Formalization

- Multiclass: Generalizes to one linear function per label:

$$g_k(\mathbf{x}) = \boldsymbol{\theta}_k^T \mathbf{x} + \theta_{0k}, \quad k = 1, 2, \dots, |\mathcal{W}|$$

- Decision rule:

$$\hat{\omega} = \arg \max_{k=1, \dots, |\mathcal{W}|} g_k(\mathbf{x}) = \arg \max_{k=1, \dots, |\mathcal{W}|} \left(\boldsymbol{\theta}_k^T \mathbf{x} + \theta_{0k} \right)$$


- Decision surface(s):

→ Each tuple $(\boldsymbol{\theta}_k^T, \theta_{0k})$ defines a hyperplane

- Decision boundaries: Pairs $\langle i, j \rangle \in \{1, \dots, |\mathcal{W}|\}$ have equal scores:

$$g_i(\mathbf{x}) = g_j(\mathbf{x}) \Rightarrow (\boldsymbol{\theta}_i - \boldsymbol{\theta}_j)^T \mathbf{x} + (\theta_{0i} - \theta_{0j}) = 0$$

Decision frontiers



Introduction

- Consider a reference set of samples $\mathcal{D} = \{(\mathbf{x}_i, \omega_i)\}_{i=1}^{|\mathcal{D}|}$

Introduction

- Consider a reference set of samples $\mathcal{D} = \{(\mathbf{x}_i, \omega_i)\}_{i=1}^{|\mathcal{D}|}$
- **Goal:** Estimate weights $\boldsymbol{\theta}$ in discriminant function $g(\mathbf{x}; \boldsymbol{\theta}) = \boldsymbol{\theta}^T \mathbf{x}$

Introduction

- Consider a reference set of samples $\mathcal{D} = \{(\mathbf{x}_i, \omega_i)\}_{i=1}^{|\mathcal{D}|}$
- **Goal:** Estimate weights θ in discriminant function $g(\mathbf{x}; \theta) = \theta^T \mathbf{x}$
→ Minimize discrepancy between predictions and ground truth (ω) on \mathcal{D} :

$$\min_{\theta} J(\theta) = \min_{\theta} \frac{1}{|\mathcal{D}|} \sum_{i=1}^{|\mathcal{D}|} \mathcal{L}(g(\mathbf{x}_i; \theta), \omega_i)$$

Introduction

- Consider a reference set of samples $\mathcal{D} = \{(\mathbf{x}_i, \omega_i)\}_{i=1}^{|\mathcal{D}|}$
- **Goal:** Estimate weights θ in discriminant function $g(\mathbf{x}; \theta) = \theta^T \mathbf{x}$
 - Minimize discrepancy between predictions and ground truth (ω) on \mathcal{D} :

$$\min_{\theta} J(\theta) = \min_{\theta} \frac{1}{|\mathcal{D}|} \sum_{i=1}^{|\mathcal{D}|} \mathcal{L}(g(\mathbf{x}_i; \theta), \omega_i)$$

- This minimization process may be solved analytically in some cases
 - Practical scenarios: iterative optimization process

Gradient descent

Gradient descent algorithm

θ_0 = arbitrary

$$\theta^{(k+1)} = \theta^{(k)} - \eta(k) \nabla J(\theta^{(k)})$$

Gradient descent

Gradient descent algorithm

θ_0 = arbitrary

$$\theta^{(k+1)} = \theta^{(k)} - \eta(k) \nabla J(\theta^{(k)})$$

- Learning rate $\eta(k)$: controls the amount of change at each iteration
→ Typical values: constant, $\eta(k) = \eta(0)/k$

Gradient descent

Gradient descent algorithm

θ_0 = arbitrary

$$\theta^{(k+1)} = \theta^{(k)} - \eta(k) \nabla J(\theta^{(k)})$$

- Learning rate $\eta(k)$: controls the amount of change at each iteration
 - Typical values: constant, $\eta(k) = \eta(0)/k$
- Process description:

Gradient descent

Gradient descent algorithm

θ_0 = arbitrary

$$\theta^{(k+1)} = \theta^{(k)} - \eta(k) \nabla J(\theta^{(k)})$$

- Learning rate $\eta(k)$: controls the amount of change at each iteration
 - Typical values: constant, $\eta(k) = \eta(0)/k$
- Process description:
 1. Criterion function $J(\theta)$ is defined

Gradient descent

Gradient descent algorithm

θ_0 = arbitrary

$$\theta^{(k+1)} = \theta^{(k)} - \eta(k) \nabla J(\theta^{(k)})$$

- Learning rate $\eta(k)$: controls the amount of change at each iteration
 - Typical values: constant, $\eta(k) = \eta(0)/k$
- Process description:
 1. Criterion function $J(\theta)$ is defined
 2. Vector θ_0 is arbitrarily chosen

Gradient descent

Gradient descent algorithm

θ_0 = arbitrary

$$\theta^{(k+1)} = \theta^{(k)} - \eta(k) \nabla J(\theta^{(k)})$$

- Learning rate $\eta(k)$: controls the amount of change at each iteration
 - Typical values: constant, $\eta(k) = \eta(0)/k$
- Process description:
 1. Criterion function $J(\theta)$ is defined
 2. Vector θ_0 is arbitrarily chosen
 3. Gradient vector $\nabla J(\theta^{(0)})$ is computed

Gradient descent

Gradient descent algorithm

θ_0 = arbitrary

$$\theta^{(k+1)} = \theta^{(k)} - \eta(k) \nabla J(\theta^{(k)})$$

- Learning rate $\eta(k)$: controls the amount of change at each iteration
 - Typical values: constant, $\eta(k) = \eta(0)/k$
- Process description:
 1. Criterion function $J(\theta)$ is defined
 2. Vector θ_0 is arbitrarily chosen
 3. Gradient vector $\nabla J(\theta^{(0)})$ is computed
 4. Value $\theta^{(1)}$: moving $\propto \eta(k)$ from $\theta^{(0)}$ in the steepest descent

Gradient descent

Gradient descent algorithm

θ_0 = arbitrary

$$\theta^{(k+1)} = \theta^{(k)} - \eta(k) \nabla J(\theta^{(k)})$$

- Learning rate $\eta(k)$: controls the amount of change at each iteration
 - Typical values: constant, $\eta(k) = \eta(0)/k$
- Process description:
 1. Criterion function $J(\theta)$ is defined
 2. Vector θ_0 is arbitrarily chosen
 3. Gradient vector $\nabla J(\theta^{(0)})$ is computed
 4. Value $\theta^{(1)}$: moving $\propto \eta(k)$ from $\theta^{(0)}$ in the steepest descent
 5. Repeat from (3) until convergence criterion is achieved

Gradient descent

Gradient descent algorithm

θ_0 = arbitrary

$$\theta^{(k+1)} = \theta^{(k)} - \eta(k) \nabla J(\theta^{(k)})$$

- Learning rate $\eta(k)$: controls the amount of change at each iteration
 - Typical values: constant, $\eta(k) = \eta(0)/k$
- Process description:
 1. Criterion function $J(\theta)$ is defined
 2. Vector θ_0 is arbitrarily chosen
 3. Gradient vector $\nabla J(\theta^{(0)})$ is computed
 4. Value $\theta^{(1)}$: moving $\propto \eta(k)$ from $\theta^{(0)}$ in the steepest descent
 5. Repeat from (3) until convergence criterion is achieved

→ Example: $|\eta(k) \nabla J(\theta^{(k)})| < \delta$

Gradient descent

Example:

- Function $J(\theta) = (\theta - 3)^2$
- Initial value $\theta_0 = -2$
- Learning rate $\eta(k) = 0.1$

Outline

① Linear models

Binary

Multiclass

Parameter estimation

② Perceptron

Introduction

Training

Limitations

Multiclass Perceptron

③ Multi-layer Perceptron

Introduction

Structure

Training

Training dynamics and regularization

Description

- Introduced by Rosenblatt in 1958

Description

- Introduced by Rosenblatt in 1958
- Practical linear **binary classifier**
 - Linear model is wrapped in a step function \Rightarrow activation


Description

- Introduced by Rosenblatt in 1958
- Practical linear **binary classifier**
 - Linear model is wrapped in a step function \Rightarrow activation
- Represents the **foundation of neural models**

Description

- Introduced by Rosenblatt in 1958
- Practical linear **binary classifier**
 - Linear model is wrapped in a step function ⇒ activation
- Represents the **foundation of neural models**


$$\hat{\omega} = \text{sign}(\theta^T \mathbf{x} + \theta_0)$$



Description

- Introduced by Rosenblatt in 1958
- Practical linear **binary classifier**
 - Linear model is wrapped in a step function ⇒ activation
- Represents the **foundation of neural models**

$$\hat{\omega} = \text{sign}(\theta^T \mathbf{x} + \theta_0) \quad \text{sign}(\cdot) = \begin{cases} +1 & \text{if } \theta^T \mathbf{x} + \theta_0 \geq 0 \\ -1 & \text{if } \theta^T \mathbf{x} + \theta_0 < 0 \end{cases}$$



Context

- Assume **binary scenario** $\mathcal{W} = \{\omega_1, \omega_2\}$

Context

- Assume **binary scenario** $\mathcal{W} = \{\omega_1, \omega_2\}$
- **Dataset** $\mathcal{D} = \{(\mathbf{x}_i, \omega_i)\}_{i=1}^{|\mathcal{D}|}$ with $\mathbf{x}_i \in \mathbb{R}^d$ and $\omega_i \in \mathcal{W}$ for $1 \leq i \leq |\mathcal{D}|$

Context

- Assume **binary scenario** $\mathcal{W} = \{\omega_1, \omega_2\}$
- **Dataset** $\mathcal{D} = \{(\mathbf{x}_i, \omega_i)\}_{i=1}^{|\mathcal{D}|}$ with $\mathbf{x}_i \in \mathbb{R}^d$ and $\omega_i \in \mathcal{W}$ for $1 \leq i \leq |\mathcal{D}|$
- Vector of **parameters** $\boldsymbol{\theta} = [\theta_0, \theta_1, \dots, \theta_d]$

Context

- Assume **binary scenario** $\mathcal{W} = \{\omega_1, \omega_2\}$
- **Dataset** $\mathcal{D} = \{(\mathbf{x}_i, \omega_i)\}_{i=1}^{|\mathcal{D}|}$ with $\mathbf{x}_i \in \mathbb{R}^d$ and $\omega_i \in \mathcal{W}$ for $1 \leq i \leq |\mathcal{D}|$
- Vector of **parameters** $\boldsymbol{\theta} = [\theta_0, \theta_1, \dots, \theta_d]$
- We may define vector $\mathbf{z} = [z_1, \dots, z_{|\mathcal{D}|}]$ such that:

$$z_i = \begin{cases} +1 & \text{if } \omega_i = \omega_1 \\ -1 & \text{if } \omega_i = \omega_2 \end{cases}$$

Context

- Assume **binary scenario** $\mathcal{W} = \{\omega_1, \omega_2\}$
- **Dataset** $\mathcal{D} = \{(\mathbf{x}_i, \omega_i)\}_{i=1}^{|\mathcal{D}|}$ with $\mathbf{x}_i \in \mathbb{R}^d$ and $\omega_i \in \mathcal{W}$ for $1 \leq i \leq |\mathcal{D}|$
- Vector of **parameters** $\boldsymbol{\theta} = [\theta_0, \theta_1, \dots, \theta_d]$
- We may define vector $\mathbf{z} = [z_1, \dots, z_{|\mathcal{D}|}]$ such that:

$$z_i = \begin{cases} +1 & \text{if } \omega_i = \omega_1 \\ -1 & \text{if } \omega_i = \omega_2 \end{cases}$$

- This way, **all samples** in \mathcal{D} will be **correctly classified** if and only if:

$$z_i \boldsymbol{\theta}^T \mathbf{x}_i > 0 \quad \forall i$$

The Perceptron Criterion

How can we estimate θ ?

The Perceptron Criterion

How can we estimate θ ? Define $J(\theta)$

The Perceptron Criterion

How can we estimate θ ? Define $J(\theta)$

- ✗ Number of misclassified samples \Rightarrow unsuitable for gradient search

The Perceptron Criterion

How can we estimate θ ? Define $J(\theta)$

- ✗ Number of misclassified samples \Rightarrow unsuitable for gradient search
- ✓ Sum of the distances to surface \mathcal{H} for misclassified samples:

$$J_p(\theta) = \sum_{i: z_i \theta^T \mathbf{x}_i \leq 0} (-z_i \theta^T \mathbf{x}_i)$$

The Perceptron Criterion

How can we estimate θ ? Define $J(\theta)$

- ✗ Number of misclassified samples \Rightarrow unsuitable for gradient search
- ✓ Sum of the distances to surface \mathcal{H} for misclassified samples:

$$J_p(\theta) = \sum_{i: z_i \theta^T \mathbf{x}_i \leq 0} (-z_i \theta^T \mathbf{x}_i)$$

→ The corresponding gradient is:

$$\nabla J_p(\theta) = \sum_{i: z_i \theta^T \mathbf{x}_i \leq 0} (-z_i \mathbf{x}_i)$$

The Perceptron Criterion

Batch Perceptron algorithm

$$\theta^{(k+1)} = \theta^{(k)} - \eta(k) \cdot \sum_{i: z_i \theta^T \mathbf{x}_i \leq 0} (-z_i \mathbf{x}_i)$$

The Perceptron Criterion

Batch Perceptron algorithm

$$\theta^{(k+1)} = \theta^{(k)} - \eta(k) \cdot \sum_{i: z_i \theta^T \mathbf{x}_i \leq 0} (-z_i \mathbf{x}_i)$$

- Instead of **batches**, correction may be done on a **singe-sample policy**:

The Perceptron Criterion

Batch Perceptron algorithm

$$\theta^{(k+1)} = \theta^{(k)} - \eta(k) \cdot \sum_{i: z_i \theta^T \mathbf{x}_i \leq 0} (-z_i \mathbf{x}_i)$$

- Instead of **batches**, correction may be done on a **singe-sample policy**:

$$\rightarrow J_p(\theta, \mathbf{x}_i) = \begin{cases} -z_i \theta^T \mathbf{x}_i & \text{if } z_i \theta^T \mathbf{x}_i \leq 0 \\ 0 & \text{otherwise} \end{cases}$$

The Perceptron Criterion

Batch Perceptron algorithm

$$\theta^{(k+1)} = \theta^{(k)} - \eta(k) \cdot \sum_{i: z_i \theta^T \mathbf{x}_i \leq 0} (-z_i \mathbf{x}_i)$$

- Instead of **batches**, correction may be done on a **singe-sample policy**:

$$\rightarrow J_p(\theta, \mathbf{x}_i) = \begin{cases} -z_i \theta^T \mathbf{x}_i & \text{if } z_i \theta^T \mathbf{x}_i \leq 0 \\ 0 & \text{otherwise} \end{cases}$$

Single-sample Perceptron algorithm

$$\theta^{(k+1)} = \theta^{(k)} + \eta(k) z_i \mathbf{x}_i$$

The Perceptron Criterion

Approach	Update frequency	Pros	Cons
Single-sample	Every sample	Fast reaction to new data stochasticity	Noisy updates
Batch	Accumulation	Smoother Stable convergence	Slower

The Perceptron Criterion

The XOR problem

The *Exclusive OR* (XOR) operation \Rightarrow **binary classification** task:

- Feature space: $\mathbf{x} \in \{0, 1\}^2$
- Label space: $\mathcal{W} = \{0, 1\}$
- Function $f_{xor}(\mathbf{x}) = \begin{cases} \omega_1 & \text{if } x_1 = x_2 \\ \omega_2 & \text{if } x_1 \neq x_2 \end{cases}$

The XOR problem

- Showcases the main **limitation** of the **perceptron** mechanism
- **Unable** to address **non-linearly separable** labels
 - Need for **non-linearities** in the scheme
- Two related **mechanisms**:
 - Stacking perceptrons into **layers**
 - Using non-linear **activation functions**