

Multiclass Perceptron

The Perceptron learner may address **multiclass scenarios** by:

Multiclass Perceptron

The Perceptron learner may address multiclass scenarios by:

1. **One-VS-Rest** scheme:

- Requires $|\mathcal{W}|$ Perceptrons
- Perceptron with the largest activation determines $\hat{\omega}$

Multiclass Perceptron

The Perceptron learner may address multiclass scenarios by:

1. **One-VS-Rest** scheme:

- Requires $|\mathcal{W}|$ Perceptrons
- Perceptron with the largest activation determines $\hat{\omega}$

2. **One-VS-One** scheme:

- Requires $|\mathcal{W}| \cdot (|\mathcal{W}| - 1)/2$ Perceptrons
- Class $\hat{\omega}$ is estimated by voting

Multiclass Perceptron

The Perceptron learner may address multiclass scenarios by:

1. **One-VS-Rest** scheme:

- Requires $|\mathcal{W}|$ Perceptrons
- Perceptron with the largest activation determines $\hat{\omega}$

2. **One-VS-One** scheme:

- Requires $|\mathcal{W}| \cdot (|\mathcal{W}| - 1)/2$ Perceptrons
- Class $\hat{\omega}$ is estimated by voting

3. **Multiple weight vectors**:

- Requires only one Perceptron
- Weight matrix with a weight vector per class: $\theta_1, \theta_2, \dots, \theta_{|\mathcal{W}|}$
- Only the weight vector associated to the class is updated while training
- Weight vector with the largest activation determines $\hat{\omega}$

Outline

① Linear models

- Binary
- Multiclass
- Parameter estimation

② Perceptron

- Introduction
- Training
- Limitations
- Multiclass Perceptron

③ Multi-layer Perceptron

- Introduction
- Structure
- Training
- Training dynamics and regularization

Introduction to MLP

- Natural evolution of the Perceptron learner

Introduction to MLP

- Natural **evolution** of the **Perceptron** learner
- Stack of **Perceptron** units \Rightarrow approximating **non-linear functions**

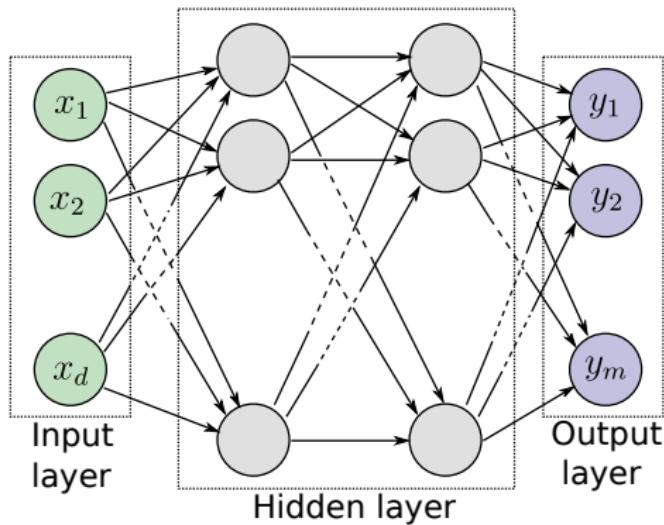
Introduction to MLP

- Natural **evolution** of the **Perceptron** learner
- **Stack of Perceptron** units \Rightarrow approximating **non-linear functions**
- Constitutes a particular case of **Feedforward Neural Networks**
 - \rightarrow Information flows in **one direction** \Rightarrow input to output

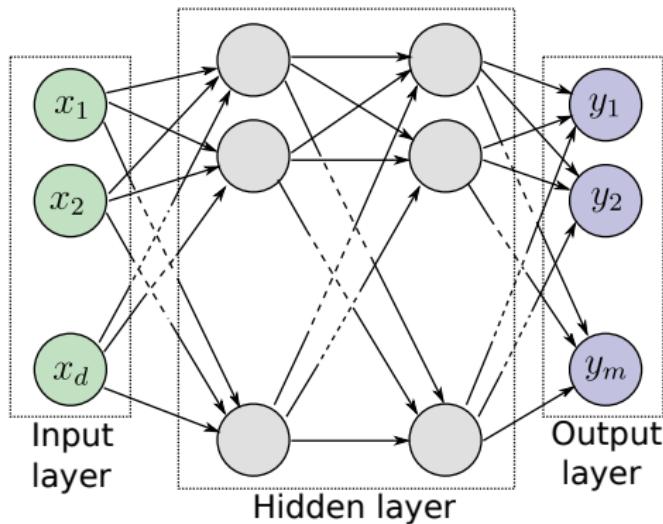
Introduction to MLP

- Natural **evolution** of the **Perceptron** learner
- Stack of **Perceptron** units \Rightarrow approximating **non-linear functions**
- Constitutes a particular case of **Feedforward Neural Networks**
 - \rightarrow Information flows in **one direction** \Rightarrow input to output
- **Universal Approximation Theorem:**
A feedforward neural network with at least one hidden layer, using a nonlinear activation, can approximate any continuous function to any desired degree of accuracy, given sufficient neurons in the hidden layer

Structure



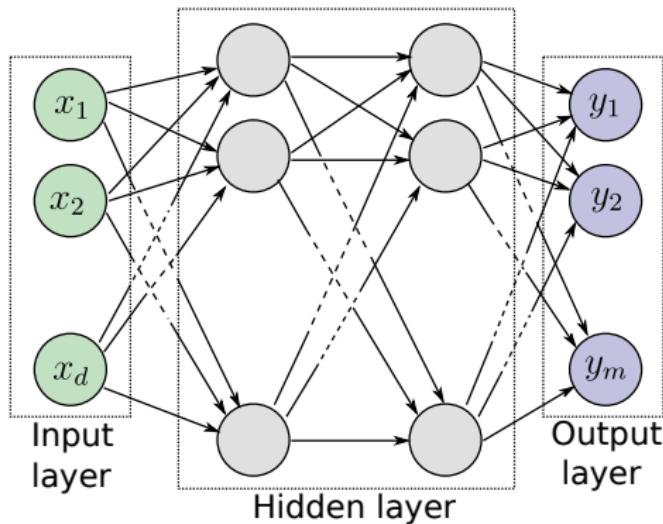
Structure



Input

- Representation space \mathbb{R}^d
⇒ d neurons

Structure



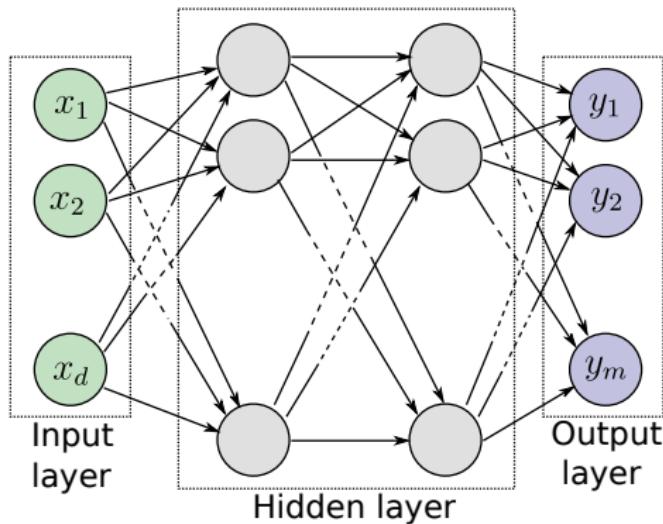
Input

- Representation space \mathbb{R}^d
⇒ d neurons

Hidden

- Non-linear activations
- Fully connected
- Configuration depends on the complexity of the function

Structure



Input

- Representation space \mathbb{R}^d
 $\Rightarrow d$ neurons

Hidden

- Non-linear activations
- Fully connected
- Configuration depends on the complexity of the function

Output

- Configuration depends on the task

Input layer

- Entrance point to the model for datum $\mathbf{x} \in \mathbb{R}^d$
- As many neurons as the dimensionality of $\mathbf{x} \Rightarrow d$ neurons
- There may be additional neurons \Rightarrow additional biases

Hidden layer

- Sequence of (vertical) stacks of Perceptron units:
 - I stacks
 - J_i neurons at the i -th stack ($1 \leq i \leq I$)

Hidden layer

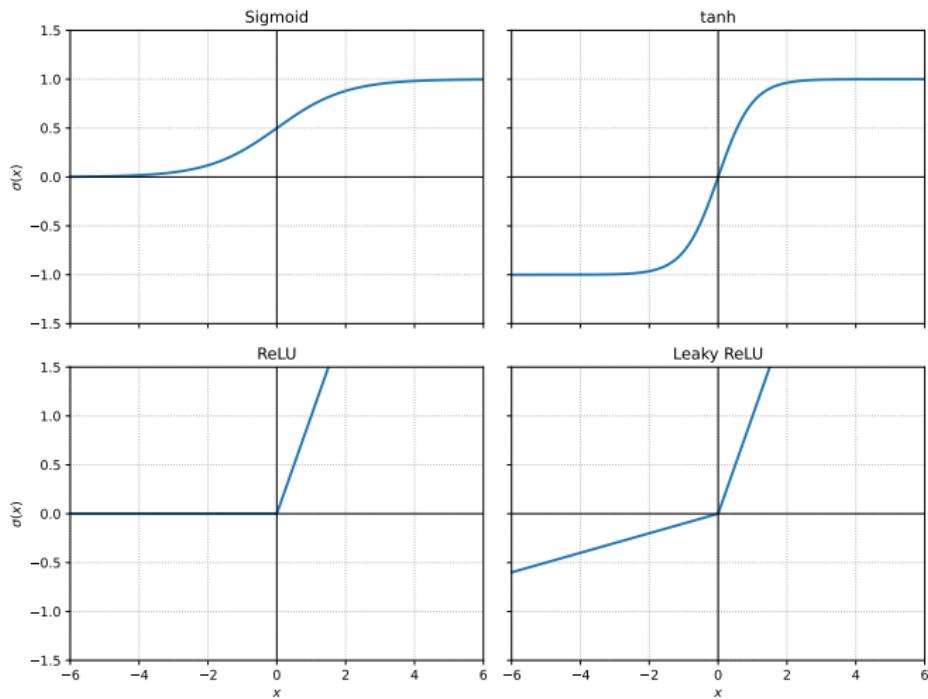
- Sequence of (vertical) stacks of Perceptron units:
 - I stacks
 - J_i neurons at the i -th stack ($1 \leq i \leq I$)
- Each stack may depict a different number of Perceptron units

Hidden layer

- Sequence of (vertical) stacks of Perceptron units:
 - I stacks
 - J_i neurons at the i -th stack ($1 \leq i \leq I$)
- Each stack may depict a different number of Perceptron units
- In general, non-linear activation functions:

$$y_{h_{i,j}} = \sigma \left(\mathbf{h}_{i,j}^T \cdot \mathbf{y}_{h_{i-1}} \right)$$

Activation functions



Activation functions

- Sigmoid:

$$\sigma(x) = \frac{e^x}{1 + e^x}$$

- Hyperbolic tangent (tanh):

$$\sigma(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

- Rectified Linear Unit (ReLU):

$$\sigma(x) = \max \{0, x\}$$

- Leaky Rectified Linear Unit (Leaky ReLU):

$$\sigma(x) = \begin{cases} x & \text{if } x \geq 0 \\ \alpha \cdot x & \text{if } x < 0 \end{cases} \quad \text{with } \alpha \in [0.1, 0.3]$$

Output layer

- Each output is a **real-valued** function: $y_j(\mathbf{x}; \boldsymbol{\theta}) \in \mathbb{R}$ with $i \leq j \leq m$

Output layer

- Each output is a **real-valued** function: $y_j(\mathbf{x}; \boldsymbol{\theta}) \in \mathbb{R}$ with $i \leq j \leq m$
- Regression tasks:
 - One **regressor per output** unit

Output layer

- Each output is a **real-valued** function: $y_j(\mathbf{x}; \boldsymbol{\theta}) \in \mathbb{R}$ with $i \leq j \leq m$
- **Regression** tasks:
 - One **regressor per output** unit
- **Classification** tasks ($|\mathcal{W}|$ classes):
 - One **unit per class** $\Rightarrow m = |\mathcal{C}|$

Output layer

- Each output is a **real-valued** function: $y_j(\mathbf{x}; \boldsymbol{\theta}) \in \mathbb{R}$ with $i \leq j \leq m$
- **Regression** tasks:
 - One **regressor per output** unit
- **Classification** tasks ($|\mathcal{W}|$ classes):
 - One **unit per class** $\Rightarrow m = |\mathcal{C}|$
 - **Highest score** is the **estimated class** $\hat{\omega} = \arg \max_{j \in \{1, \dots, m\}} y_j$

Output layer

- Each output is a **real-valued** function: $y_j(\mathbf{x}; \boldsymbol{\theta}) \in \mathbb{R}$ with $i \leq j \leq m$
- **Regression** tasks:
 - One **regressor per output** unit
- **Classification** tasks ($|\mathcal{W}|$ classes):
 - One **unit per class** $\Rightarrow m = |\mathcal{C}|$
 - **Highest score** is the **estimated class** $\hat{\omega} = \arg \max_{j \in \{1, \dots, m\}} y_j$
 - **Softmax activation as estimated probability**:

$$\hat{P}(\omega | \mathbf{x}) = \frac{e^{y_\omega}}{\sum_{j=1}^{|\mathcal{W}|} e^{y_j}} \in [0, 1]$$

Output layer

- Each output is a **real-valued** function: $y_j(\mathbf{x}; \boldsymbol{\theta}) \in \mathbb{R}$ with $i \leq j \leq m$

- **Regression** tasks:
 - One **regressor per output** unit

- **Classification** tasks ($|\mathcal{W}|$ classes):
 - One **unit per class** $\Rightarrow m = |\mathcal{C}|$
 - **Highest score** is the **estimated class** $\hat{\omega} = \arg \max_{j \in \{1, \dots, m\}} y_j$
 - **Softmax activation as estimated probability**:

$$\hat{P}(\omega | \mathbf{x}) = \frac{e^{y_\omega}}{\sum_{j=1}^{|\mathcal{W}|} e^{y_j}} \in [0, 1]$$

→ **Binary** case ($|\mathcal{C}| = 2$) may be modeled with a **single neuron**:

$$\hat{\omega} = \begin{cases} \omega_1 & \text{if } y \geq 0 \\ \omega_2 & \text{if } y < 0 \end{cases}$$

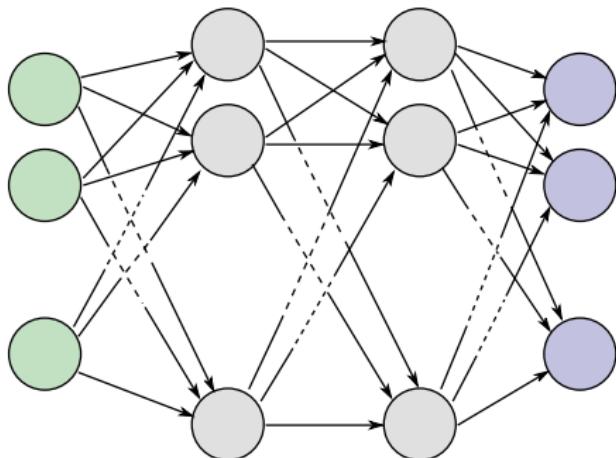
Contextualization

- Training is done following the *forward-backward* propagation

Contextualization

- Training is done following the *forward-backward* propagation

- Steps:

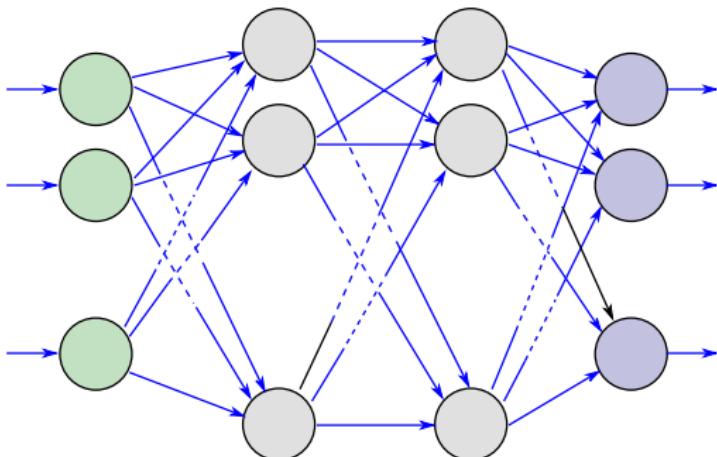


Contextualization

- Training is done following the *forward-backward* propagation

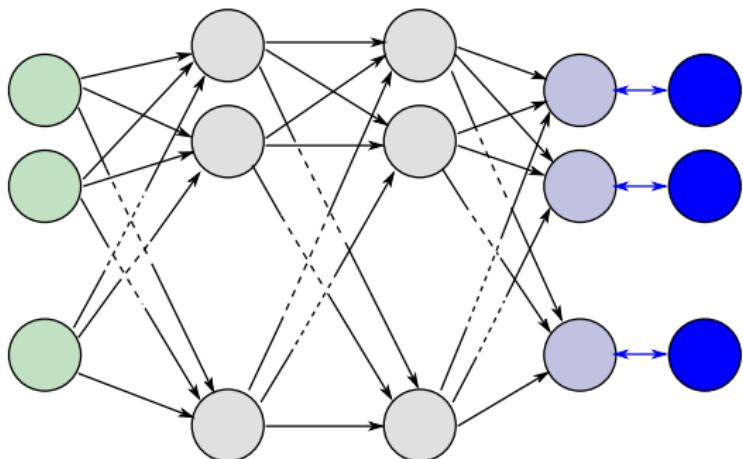
- Steps:

1. Forward pass



Contextualization

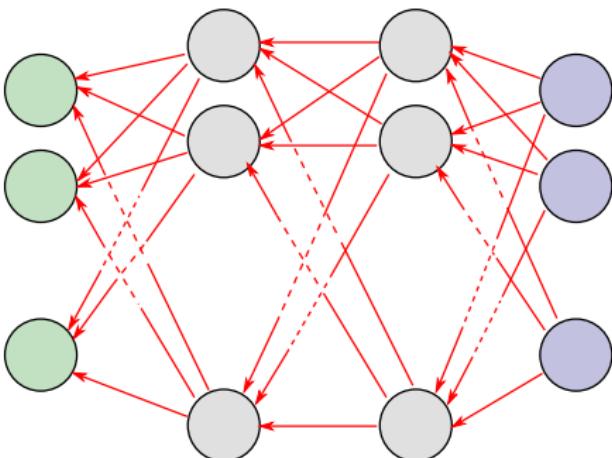
- Training is done following the *forward-backward* propagation
- Steps:
 1. Forward pass
 2. Loss computation



Contextualization

- Training is done following the *forward-backward* propagation

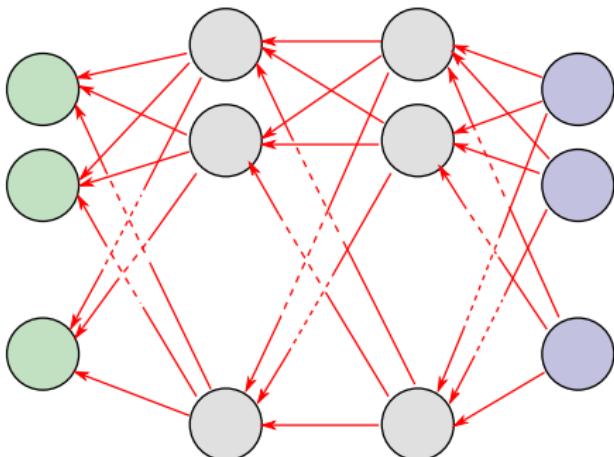
- Steps:
 1. Forward pass
 2. Loss computation
 3. Backward pass
 4. Weight update



Contextualization

- Training is done following the *forward-backward* propagation

- Steps:
 1. Forward pass
 2. Loss computation
 3. Backward pass
 4. Weight update

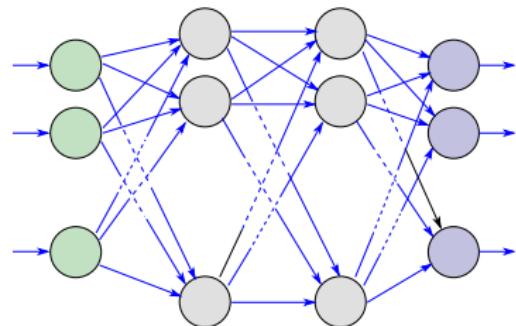


- The process is done until a convergence criterion is reached

Forward-backward training

1. Forward pass:

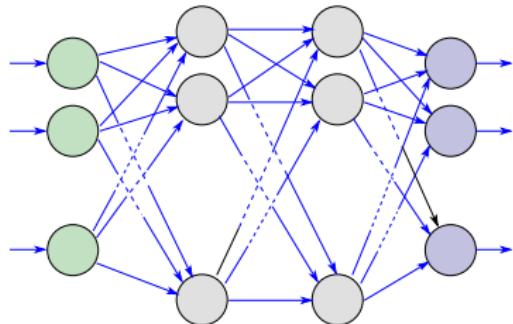
- Datum x goes from **input to output**
- Weights θ modify x across the network



Forward-backward training

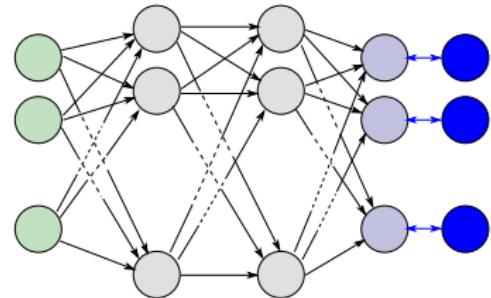
1. Forward pass:

- Datum x goes from **input to output**
- Weights θ modify x across the network



2. Loss computation:

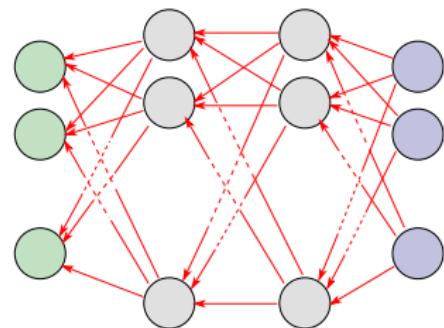
- Estimation \hat{y} is compared to reference y
- Depends on the task:
 - Regression: MAE, MSE
 - Classification: Categorical Cross Entropy



Forward-backward training

3. Backward pass (backpropagation) :

- Gradient of the loss function with respect to θ
- Chain rule for the derivatives of the (nested) equation
- Gradients are propagated from output to input



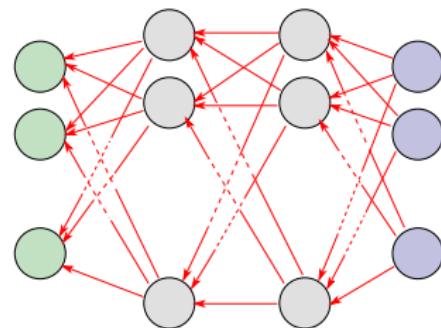
Forward-backward training

3. Backward pass (backpropagation) :

- Gradient of the loss function with respect to θ
- Chain rule for the derivatives of the (nested) equation
- Gradients are propagated from output to input

4. Weight update:

- Update parameters θ (weights and biases)
- Optimization methods:
 - Stochastic Gradient Descent
 - Momentum, RMSProp, Adam, ...



Classification scenarios

- Number of neurons must match labels $\Rightarrow m = |\mathcal{W}|$

Classification scenarios

- Number of **neurons** must match **labels** $\Rightarrow m = |\mathcal{W}|$
- Consider a case of $m = 3$:

Classification scenarios

- Number of **neurons** must match **labels** $\Rightarrow m = |\mathcal{W}|$
- Consider a case of $m = 3$:

Ideal case

Class A: [1, 0, 0]

Class B: [0, 1, 0]

Class C: [0, 0, 1]

Classification scenarios

- Number of **neurons** must match **labels** $\Rightarrow m = |\mathcal{W}|$
- Consider a case of $m = 3$:

Ideal case

Class A: [1, 0, 0]
Class B: [0, 1, 0]
Class C: [0, 0, 1]

Real case

Class A: [0.67, 0.10, 0.23]
Class B: [0.05, 0.80, 0.15]
Class C: [0.25, 0.25, 0.50]

Classification scenarios

- Number of **neurons** must match **labels** $\Rightarrow m = |\mathcal{W}|$
- Consider a case of $m = 3$:

Ideal case

Class A: [1, 0, 0]
Class B: [0, 1, 0]
Class C: [0, 0, 1]

Real case

Class A: [0.67, 0.10, 0.23]
Class B: [0.05, 0.80, 0.15]
Class C: [0.25, 0.25, 0.50]

- **Training** \Rightarrow reference data must be adequately \Rightarrow **one-hot encoding**
 \rightarrow **Mapping** function: $[\omega_1, \dots, \omega_{|\mathcal{W}|}] \rightarrow [0, 1]^{|\mathcal{W}|}$

Classification scenarios

- Number of **neurons** must match **labels** $\Rightarrow m = |\mathcal{W}|$
- Consider a case of $m = 3$:

Ideal case

Class A: [1, 0, 0]
 Class B: [0, 1, 0]
 Class C: [0, 0, 1]

Real case

Class A: [0.67, 0.10, 0.23]
 Class B: [0.05, 0.80, 0.15]
 Class C: [0.25, 0.25, 0.50]

- **Training** \Rightarrow reference data must be adequately \Rightarrow **one-hot encoding**
 \rightarrow **Mapping** function: $[\omega_1, \dots, \omega_{|\mathcal{W}|}] \rightarrow [0, 1]^{|\mathcal{W}|}$
- **Categorical Cross Entropy** \Rightarrow Typical **loss** function for classification

$$\mathcal{L}(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_{j=1}^{|\mathcal{W}|} y_j \cdot \log(\hat{y}_j)$$

Classification scheme

Consider the following [example](#):

Classification scheme

Consider the following [example](#):

- Feature space \mathbb{R}^d with $d = 2$ dimensions

Classification scheme

Consider the following example:

- Feature space \mathbb{R}^d with $d = 2$ dimensions
- Three possible categories: $\mathcal{W} = \{\omega_1, \omega_2, \omega_3\}$

Classification scheme

Consider the following example:

- Feature space \mathbb{R}^d with $d = 2$ dimensions
- Three possible categories: $\mathcal{W} = \{\omega_1, \omega_2, \omega_3\}$
- Test element $\mathbf{x} = [0.4, -2.2]$ with $\omega = \omega_2$

Classification scheme

Consider the following example:

- Feature space \mathbb{R}^d with $d = 2$ dimensions
- Three possible categories: $\mathcal{W} = \{\omega_1, \omega_2, \omega_3\}$
- Test element $\mathbf{x} = [0.4, -2.2]$ with $\omega = \omega_2 \Rightarrow \mathbf{y} = [0, 1, 0]$

Classification scheme

Consider the following example:

- Feature space \mathbb{R}^d with $d = 2$ dimensions
- Three possible categories: $\mathcal{W} = \{\omega_1, \omega_2, \omega_3\}$
- Test element $\mathbf{x} = [0.4, -2.2]$ with $\omega = \omega_2 \Rightarrow \mathbf{y} = [0, 1, 0]$
- Predicted vector $\hat{\mathbf{y}} = [0.1, 0.7, 0.2]$

Classification scheme

Consider the following example:

- Feature space \mathbb{R}^d with $d = 2$ dimensions
- Three possible categories: $\mathcal{W} = \{\omega_1, \omega_2, \omega_3\}$
- Test element $\mathbf{x} = [0.4, -2.2]$ with $\omega = \omega_2 \Rightarrow \mathbf{y} = [0, 1, 0]$
- Predicted vector $\hat{\mathbf{y}} = [0.1, 0.7, 0.2]$

The associated loss is:

$$\begin{aligned}\mathcal{L}(\mathbf{y}, \hat{\mathbf{y}}) &= - \sum_{j=1}^3 y_j \cdot \log(\hat{y}_j) \\ &= -[0 \cdot \log(0.1) + 1 \cdot \log(0.7) + 0 \cdot \log(0.2)] = 0.357\end{aligned}$$

Practical considerations

MLPs (and other Neural models) depict remarkable limitations:

Practical considerations

MLPs (and other Neural models) depict remarkable limitations:

1. **Overfitting:** *Memorizing instead of learning*

Practical considerations

MLPs (and other Neural models) depict remarkable limitations:

1. **Overfitting:** *Memorizing* instead of *learning*
2. **Vanishing/exploding gradients:** Gradients become negligible/excessively large producing stagnation/oscillations

Practical considerations

To prevent overfitting \Rightarrow regularization strategies:

Practical considerations

To prevent **overfitting** \Rightarrow **regularization** strategies:

- **L1/L2 regularization:** Adding additional penalties during training

Practical considerations

To prevent **overfitting** \Rightarrow **regularization** strategies:

- **L1/L2 regularization:** Adding additional penalties during training
- **Dropout:** Randomly disconnecting neurons during training

Practical considerations

To prevent **overfitting** \Rightarrow **regularization** strategies:

- **L1/L2 regularization**: Adding additional penalties during training
- **Dropout**: Randomly disconnecting neurons during training
- **Early stopping**: Stop training as validation loss increases

Practical considerations

To prevent **overfitting** \Rightarrow **regularization** strategies:

- **L1/L2 regularization**: Adding additional penalties during training
- **Dropout**: Randomly disconnecting neurons during training
- **Early stopping**: Stop training as validation loss increases
- **Data augmentation**: Apply transformations on (train) data to artificially expand the size of the set

Practical considerations

To prevent vanishing/exploding gradients:

Practical considerations

To prevent vanishing/exploding gradients:

- **Non-linear activations:** (Leaky) ReLU palliate vanishing issue

Practical considerations

To prevent vanishing/exploding gradients:

- **Non-linear activations:** (Leaky) ReLU palliate vanishing issue
- **Initialization strategies:** Weight initialization policies

Practical considerations

To prevent **vanishing/exploding** gradients:

- **Non-linear activations:** (Leaky) ReLU palliate vanishing issue
- **Initialization strategies:** Weight initialization policies
- **Batch/Layer normalization:** Maintaining activations in stable ranges

Practical considerations

To prevent **vanishing/exploding** gradients:

- **Non-linear activations:** (Leaky) ReLU palliate vanishing issue
- **Initialization strategies:** Weight initialization policies
- **Batch/Layer normalization:** Maintaining activations in stable ranges
- **Gradient clipping:** Limiting gradient value above threshold

T5: Linear methods and Perceptron

Fundamentos del Aprendizaje Automático

Curso 2025/2026