





Redes Neuronales y Aprendizaje Profundo



Redes Neuronales y Aprendizaje Profundo

Bloque 2: Convolutional Neural Networks

Fundamentals



Convolutional Neural Networks

- **Objectives**
 - **Explain the limitations** of fully connected neural networks when applied to image data **and justify the need for convolutional architectures.**
 - **Describe the convolution operation** in convolutional neural networks, including the role of kernels, feature maps, stride, and padding.
 - **Analyze the effect of basic architectural choices** (e.g., number of filters, kernel size, pooling) on the dimensionality of feature maps.
 - **Interpret how spatial structure and local connectivity are exploited by CNNs** to learn meaningful visual representations.

Convolutional Neural Networks

- **Objectives**
 - **Identify and explain the main building blocks of a CNN**, including convolutional layers, activation functions, pooling layers, and fully connected layers.
 - **Design a simple CNN architecture** for an image classification task at a conceptual level.

Contents

- Motivation
- The Convolution Operation
- Core Building Blocks of a CNN
- Hierarchical Feature Learning
- Designing a Basic CNN Architecture

Motivation


- Convolutional Neural Networks (CNNs) are specifically tailored for computer vision tasks:

Classification



Dog, Dog

Detection



Segmentation




Synthesis





Motivation

- How we can use a FCN to classify images?





Motivation

- How we can use a FCN to classify images?




Motivation

- How many **parameters** do we have to learn?




Motivation

- And now?



1024x1024x3



$X \rightarrow 3M$

First layer

$W^1 \rightarrow 3M \times 1000$

$b^1 \rightarrow 1000$

Second layer

$W^2 \rightarrow 1000 \times 100$

$b^2 \rightarrow 100$

Output layer


$W^3 \rightarrow 100 \times 1$

$b^3 \rightarrow 1$


$$\text{Total} = 3.145.729.000 + 100.100 + 101 = \mathbf{3.145.829.201}$$

Motivation

- And now?




1024x1024x3



$$\text{Total} = \underline{3.145.729.000} + 100.100 + 101 = \underline{\underline{3.145.829.201}}$$


Motivation

- **Parameter Explosion**
 - Parameters grow linearly with input size
 - High risk of overfitting
 - Computationally inefficient
- **Fully connected networks do not scale to realistic image sizes**



Motivation


- **Images are not just vectors**
 - An image is a 3D tensor: $H \times W \times C$
 - Strong spatial correlations between neighboring pixels
 - Local patterns (edges, corners, textures) matter more than absolute pixel values



	R	5	0	0	60	355
G	0.10	0.29	4.09	8.25	4.58	4.85
B	(.25)	(.90)	7.30	1.85	2.35	355
(.95)	6.20	(1620)	256	2.29	255	380)
(.92)	1.30	0.59	379	2.55	355	2.90
0.95	1.22	2.88	325	2.57	235	8.30)
(.26)	9.30	1.20	258	3.28	285	1.27
0.27	9.00	1323	2.56	(1.36)	255	(1.56)
(.24)	1.70	345	1635	325	388	(50)
(2.29)	635	285	1025	6.29	(36	(25


Motivation

- **Loss of spatial information.** Flattening removes:
 - Local neighborhoods




Motivation

- **Loss of spatial information.** Flattening removes:
 - Local neighborhoods
 - Relative positions: vision relies on spatial relationships: edges, corners, and object parts are defined by their relative positions.




Motivation

- **Loss of spatial information.** Flattening removes:
 - Local neighborhoods
 - Relative positions
 - Translation awareness: network doesn't understand they are the same pattern but traslated




Motivation

- If patterns are local and can appear anywhere in the image → we should **learn them locally and reuse them everywhere**.
- **Design requirements**
 - Local connectivity
 - Weight sharing
 - Translation invariance
 - Efficient parameterization



Convolutional operation

- **Parameters**
 - Filter size
 - Stride
 - Padding



Convolutional operation. Parameters

- **Filter size**


- **Local Receptive Fields**

- Each neuron sees only a **small region of the image**
 - Not the whole image
 - Only a local patch
 - Same size as the kernel




3x3

- Used in most modern CNNs
- VGG, ResNet...
- 9 params




5x5

- Early CNNs
- Inception
- 25 params.



7x7



- Modern lightweight networks avoid large kernels
- First layer ResNet
- 49 params



Convolutional operation. Parameters

- **Stride**

- Defines the **step size** of the kernel as it moves across the image.
- Stride = 1 → the filter moves one pixel at a time.
- Stride > 1 → the filter skips positions.
- Larger stride → smaller output feature map.
- Increasing stride reduces spatial resolution and computational cost.
- Most convolutions use stride = 1.
- Stride = 2 is typically used for spatial downsampling.
- Stride > 2 → rare




Convolutional operation. Parameters

- **Padding**

- Adding borders to the image
- Prevents shrinking
- Preserves spatial dimensions
- Helps detect patterns at borders
- Types:
 - Valid (no padding)
 - Same (keeps size)

- **Example**

- What is the output size using a filter of 3×3 and stride 1 for an image of 32×32 ?
 - Without padding
 - With padding = 1



Convolutional operation. Parameters

- **Exercise:** Calculate the feature map for the following parameters
 - Image 64x64
 - Filter 3x3
 - Stride 1
 - Padding 0
 - Image 64x64
 - Filter 3x3
 - Stride 1
 - Padding 1
 - Image 64x64
 - Filter 3x3
 - Stride 2
 - Padding 0
 - Image 64x64
 - Filter 5x5
 - Stride 2
 - Padding 2

Convolutional operation. Parameters

- **Exercise:** Calculate the feature map for the following parameters

- Image 64x64
 - Filter 3x3
 - Stride 1
 - Padding 0
- **62 x 62**

- Image 64x64
 - Filter 3x3
 - Stride 1
 - Padding 1
- **64 x 64**

- Image 64x64
 - Filter 3x3
 - Stride 2
 - Padding 0
- **31 x 31**

- Image 64x64
 - Filter 5x5
 - Stride 2
 - Padding 2
- **32 x 32**

Stride controls downsampling, while padding controls size preservation.

Convolutional operation. Parameters

- **Size of the feature map**

- The result of a convolution filter with size ($f \times f$) to an image of (h, w) size with a padding p is:


$$(h + 2p - f + 1, w + 2p - f + 1)$$

- The result of a convolution filter with size $(f \cdot f)$ to an image of (h, w) size with a padding p and a stride of s is:

$$\left(\left\lfloor \frac{h + 2p - f}{s} \right\rfloor + 1, \left\lfloor \frac{w + 2p - f}{s} \right\rfloor + 1 \right)$$


Convolutional operation. Volumes

- Convolution over volumes. Each filter:
 - Spans all input channels
 - Produces one feature map



Input ($H \times W \times 3$)


Kernel ($3 \times 3 \times 3$)




Convolutional operation

- **Multiple filters**

- Each filter learns different things (vertical and horizontal borders, corners,...)
- Filters compete to specialize



Input ($H \times W \times 3$) N Kernels ($3 \times 3 \times 3$)



$$\text{N Feature maps } \left(\left(\left\lfloor \frac{H+2p-f}{s} \right\rfloor + 1, \left\lfloor \frac{W+2p-f}{s} \right\rfloor + 1 \right) \right)$$


Convolutional building blocks

- All together. Including bias

Input ($H \times W$)




Kernel (3×3)



b


$$W^{(L)} \quad | \quad b^{(L)}$$

Feature map



Activation

σ



$$a^{(L)} = \sigma(z^{(L)})$$

$$a^{(L-1)} = X$$

$$z^{(L)} = W^{(L)} * a^{(L-1)} + b^{(L)}$$

* Convolution

$$z_{i,j,k}^{(L)} = \sum_{m,n,c} W_{m,n,c,k}^{(L)} a_{i+m,j+n,c}^{(L-1)} + b_k^{(L)}$$


Convolutional building blocks

- Number of **parameters to learn** in a Convolutional layer
 - Let c_{L-1} the number of channels of the previous layer L of a convolutional layer
 - f the filter height and width and
 - c the number of filters in the layer.


$$\# \text{ params} = (f \cdot f \cdot c_{L-1} + 1) \cdot c_L$$

Convolutional building blocks

- And now?




$1024 \times 1024 \times 3$



$$\text{Total} = \underline{3.145.729.000} + 100.100 + 101 = \underline{\underline{3.145.829.201}}$$


Convolutional building blocks

- And now?




1024x1024x3

10 filters



X W^1, b^1



$X \rightarrow 3M$

First layer

$W^1 \rightarrow 270$

$b^1 \rightarrow 10$

Second layer

$W^2 \rightarrow \dots$

$b^2 \rightarrow \dots$

Output layer

$W^3 \rightarrow \dots$

$b^3 \rightarrow \dots$


Difference = 3.145.729.000 vs 280 parameters

Convolutional building blocks

- And now?

10 filters

$X \rightarrow 3M$




Difference = 3.145.729.000 vs 280 parameters

Convolutional building blocks

- **Exercise**

- Calculate the number of parameters




Convolutional building blocks

- **Pooling layers:** Reduces spatial dimensions
 - Reduce computation
 - Reduce overfitting
 - Increase receptive field
 - Introduce local translation invariance

One Feature Map


2	3	2	0
5	-2	2	8
-1	-6	7	3
-4	-5	4	2



(a) Max-Pooling

One Feature Map

2	3	2	0
5	-2	2	8
-1	-6	7	3
-4	-5	4	2




(b) Average-Pooling

No parameters to learn !!

Hierarchical Feature Learning

- CNNs learn representations in stages:
 - Early layers → low-level features
 - Middle layers → intermediate patterns
 - Deep layers → high-level concepts




Hierarchical Feature Learning

- CNNs learn representations in stages:
 - Early layers → low-level features
 - Middle layers → intermediate patterns
 - Deep layers → high-level concepts
- Typically learn:
 - Edges (vertical, horizontal)
 - Corners
 - Color contrasts
 - Simple textures




Hierarchical Feature Learning

- CNNs learn representations in stages:
 - Early layers → low-level features
 - **Middle layers** → intermediate patterns
 - Deep layers → high-level concepts
- Combine low-level features into:
 - Textures
 - Repeated patterns
 - Parts of objects (eyes, wheels,...).




Hierarchical Feature Learning

- CNNs learn representations in stages:
 - Early layers → low-level features
 - Middle layers → intermediate patterns
 - **Deep layers** → high-level concepts



Designing a basic CNN architecture

- A typical CNN has:
 - Several **convolution plus pooling layers**
 - Each responsible for feature extraction at different levels of abstraction
 - A **final classifier**






As we go deeper:

- Size ↓
- Channels ↑
- Spatial resolution ↓
- Semantic richness ↑

Exercise

- Calculate the activation size and the number of parameters of LeNet



Exercise (optional)

- Calculate the Forward pass and Backward pass for the first convolutional layer

